


Deep Learning Final Exam

...

Kaggle : Image Generation Model

- Goal is to build a GAN that can generate/discriminate images from e-commerce clothing store.
- Dataset : fashion MNIST
- Challenges :
 - Limited amount of computing resources

What's the data look like ?



Pre-processing

```
def scale_images(data):  
    image = data['image']  
    return image / 255  
  
dataset = tfds.load('fashion_mnist', split='train')  
dataset = dataset.map(scale_images) # Parallelizing Data Transformation  
dataset = dataset.cache()  
dataset = dataset.shuffle(10000) # Adjusted buffer size for memory management  
dataset = dataset.batch(256) # Adjust based on your GPU  
dataset = dataset.prefetch(tf.data.experimental.AUTOTUNE) # Auto-tune the prefetch size  
dataset = strategy.experimental_distribute_dataset(dataset) # Leverage the two GPU
```

Hardware challenges

- Started with google Colab but quickly limited
- Managed to have access to NVIDIA Tesla P4 GPU but the cooling fan was not strong enough to prevent overheating during training
- Found out that Kaggle allow for 30hrs of usage of 2 T4 freely

Training challenge

- Not owning the hardware and large training time make it more difficult perform some trial & errors approach
- Even after 12 hrs of training our 1st model, the results were not satisfying
- Decided to change the approach and to finetune an existing model (ResNet50) instead of creating it from scratch

Initial generator


```
def build_generator():  
    model = Sequential()  
  
    # Input layer  
    model.add(Dense(7*7*128, input_dim=128))  
    model.add(LeakyReLU(0.2))  
    model.add(Reshape((7,7,128)))  
  
    # Upsampling block 1  
    model.add(UpSampling2D())  
    model.add(Conv2D(128, 5, padding='same'))  
    model.add(BatchNormalization())  
    model.add(LeakyReLU(0.2))  
  
    # Upsampling block 2  
    model.add(UpSampling2D())  
    model.add(Conv2D(128, 5, padding='same'))  
    model.add(BatchNormalization())  
    model.add(LeakyReLU(0.2))  
  
    # Convolutional block 1  
    model.add(Conv2D(128, 4, padding='same'))  
    model.add(BatchNormalization())  
    model.add(LeakyReLU(0.2))  
  
    # Convolutional block 2  
    model.add(Conv2D(128, 4, padding='same'))  
    model.add(BatchNormalization())  
    model.add(LeakyReLU(0.2))  
  
    # Output conv layer  
    model.add(Conv2D(1, 4, padding='same', activation='tanh'))  
  
    return model
```

Total params: 2,157,185 (8.23 MB)

Trainable params: 2,156,161 (8.23 MB)

Non-trainable params: 1,024 (4.00 KB)

Generator output (no training)



Initial discriminator


```
def build_discriminator():  
    model = Sequential()  
  
    # First Conv Block  
    model.add(Conv2D(32, 5, strides=(2, 2), padding='same', input_shape=(28, 28, 1)))  
    model.add(LeakyReLU(0.2))  
  
    # Second Conv Block  
    model.add(Conv2D(64, 5, strides=(2, 2), padding='same'))  
    model.add(LeakyReLU(0.2))  
  
    # Third Conv Block  
    model.add(Conv2D(128, 5, strides=(2, 2), padding='same'))  
    model.add(LeakyReLU(0.2))  
  
    # Fourth Conv Block  
    model.add(Conv2D(256, 5, strides=(2, 2), padding='same'))  
    model.add(LeakyReLU(0.2))  
  
    # Flatten then pass to dense layer  
    model.add(Flatten())  
    model.add(Dense(1, activation='sigmoid'))  
  
    return model  
  
discriminator = build_discriminator()  
discriminator.summary()
```

Total params: 1,077,505 (4.11 MB)

Trainable params: 1,077,505 (4.11 MB)

Non-trainable params: 0 (0.00 B)

Final image generation after ~18hrs of training




Second generator

```
def build_generator():  
    model = tf.keras.Sequential([  
        tf.keras.layers.Input(shape=(100,)),  
        tf.keras.layers.Dense(7*7*256, use_bias=False),  
        tf.keras.layers.BatchNormalization(),  
        tf.keras.layers.LeakyReLU(),  
        tf.keras.layers.Reshape((7, 7, 256)),  
        tf.keras.layers.Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same', use_bias=False),  
        tf.keras.layers.BatchNormalization(),  
        tf.keras.layers.LeakyReLU(),  
        tf.keras.layers.UpSampling2D(),  
        tf.keras.layers.Conv2DTranspose(64, (5, 5), strides=(1, 1), padding='same', use_bias=False),  
        tf.keras.layers.BatchNormalization(),  
        tf.keras.layers.LeakyReLU(),  
        tf.keras.layers.UpSampling2D(),  
        tf.keras.layers.Conv2DTranspose(1, (5, 5), strides=(1, 1), padding='same', use_bias=False, activation='tanh')  
    ])  
    return model
```

Second discriminator

```
def build_discriminator():  
    # Define the input shape and preprocess inputs  
    inputs = Input(shape=(28, 28, 1))  
    x = UpSampling2D(size=(8, 8))(inputs) # Upsample to 224x224  
    x = Conv2D(3, (3, 3), padding='same', activation='relu')(x) # Convert to 3 channels  
  
    # Utilize ResNet50 as a feature extractor  
    resnet_model = ResNet50(include_top=False, input_shape=(224, 224, 3), pooling='avg')  
    resnet_model.trainable = False # Freeze the model  
    x = resnet_model(x)  
  
    # Flatten the output and add a Dense layer for binary classification  
    x = Flatten()(x)  
    outputs = Dense(1, activation='sigmoid')(x)  
  
    # Create the model  
    model = Model(inputs, outputs)  
    return model
```

Results after ~4hrs of training



Conclusion

- Hardware/Computing power and training time are a major hurdle in training and deploying deep learning model

Areas of improvement

- Go back to initial architecture without using pre-trained model
- Setup a more powerful NN training lab
- Give more training time to the model
- Experiment with different architecture for generator and discriminator